



**WAITING FOR STUDENTS TO TURN VIDEOS ON SO  
I DON'T FEEL LIKE I'M TALKING TO AN EMPTY ROOM**

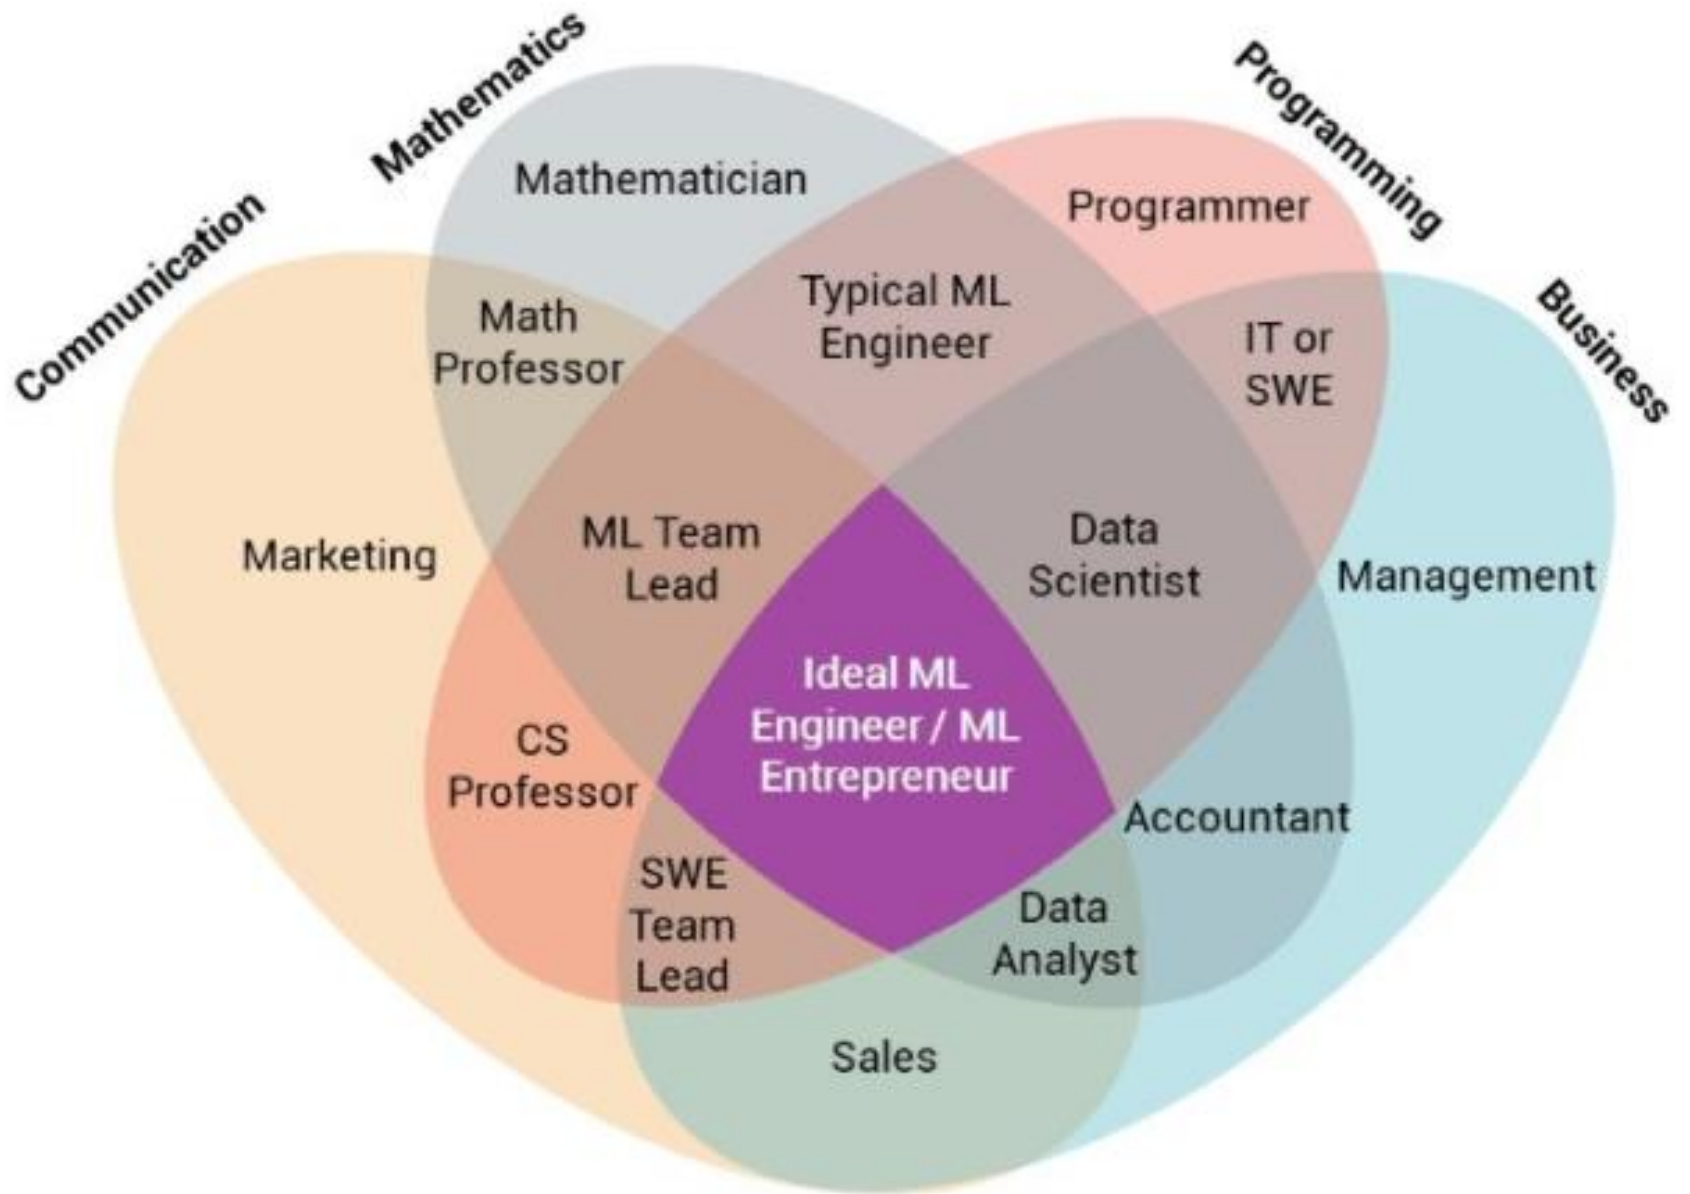



# THI GIÁC MÁY TÍNH TRONG TƯƠNG TÁC NGƯỜI MÁY

Computer Vision in Human - Computer Interaction



**AI Career Pathways**






# Executive Summary

Developing an AI project development life cycle involves five distinct tasks:

- **Data engineering:** People responsible for data engineering prepare data and transform data into formats that other team members can use.
- **Modeling:** People assigned to modeling look for patterns in data that can help a company predict outcomes of various decisions, identify business risks and opportunities, or determine cause-and-effect relationships.
- **Deployment:** People in charge of deployment take a stream of data, combine it with a model, and test the integration before putting the model into production.
- **Business analysis:** Team members responsible for business analysis evaluate a deployed model's performance and business value and adjust accordingly to maximize benefit or abandon unproductive models.
- **AI infrastructure:** People who work in AI infrastructure build and maintain reliable, fast, secure, and scalable software systems to help people working in data engineering, modeling, deployment and business analysis.





# Data engineering work includes:

| Subtask                      | Examples  | Skills involved   |
|------------------------------|---|---|
| Defining data requirements   | <ul style="list-style-type: none"><li>- Creating a data model</li><li>- Defining the features of high-quality data</li><li>- Defining the covariates to be collected to achieve a desired functionality</li><li>- Providing feedback regarding the clarity and completeness of data requirements</li></ul>  | Machine learning<br>Business acumen<br>Software engineering |
| Collecting data              | <ul style="list-style-type: none"><li>- Setting up a Mechanical Turk project</li><li>- Collecting data by manually taking images of cats</li><li>- Coding a JavaScript tracker on a website to collect user data</li><li>- Scraping the web and, if necessary, synchronizing data from different sources</li></ul>  | Machine learning<br>Software engineering                    |
| Labeling data                | <ul style="list-style-type: none"><li>- Drawing bounding boxes on images</li><li>- Building a labeling pipeline using Mechanical Turk</li><li>- Writing a labeling tutorial for workers</li><li>- Relabeling mislabeled data</li><li>- Evaluate the labeling performance of workers</li></ul>   | Machine learning<br>Software engineering                    |
| Inspecting and cleaning data | <ul style="list-style-type: none"><li>- Replacing all non-usable structured data records by NaN using a Python library (e.g. pandas)</li><li>- Converting a continuous feature into a categorical feature using bucketing</li><li>- Reformatting a data set (for instance, converting images to jpeg and squaring them)</li><li>- Cleaning a text dataset (for instance, removing special characters)</li></ul> | Machine learning<br>Algorithmic coding                      |

# Data engineering work includes :

---

|   |   |  |
|---|---|--|
| Augmenting data                         | <ul style="list-style-type: none"><li>- Writing a Python script using skimage to rotate, warp, translate, or blur images</li><li>- Using test-time augmentation to reduce the variance of an algorithm</li><li>- Synthesizing speech by overlaying distinct audio signals</li></ul>                                 | Machine learning<br>Algorithmic coding               |
| Moving data and building data pipelines | <ul style="list-style-type: none"><li>- Writing a script to allow online learning for a model</li><li>- Designing an ETL system</li><li>- Writing a script to preprocess training data and send it as input to a model automatically</li><li>- Writing a script to record model predictions in a database</li></ul> | Domain-specific (for instance, data query) languages |
| Querying data                           | <ul style="list-style-type: none"><li>- Pulling data from a database</li></ul>  | Domain-specific (for instance, data query) languages |
| Tracing data                            | <ul style="list-style-type: none"><li>- Keeping track of data sources</li><li>- Setting up a data version control system</li></ul>  | Software engineering                                 |

# Modeling work includes :

| Subtask  | Examples   | Skills involved   |
|--|--|---|
| Training machine learning models   | <ul style="list-style-type: none"> <li>- Using one of the following methods: Linear Regression, Logistic Regression, Decision Trees, Random Forest, XGBoost, Support Vector Machines, K-means, K-Nearest Neighbors, Neural Networks, Principal Component Analysis, Naive Bayes Classifier, Lasso/Ridge regression, etc.</li> </ul> | Machine learning<br>Algorithmic coding<br>Mathematics<br>Data science |
| Fitting probabilistic or statistical models                                | <ul style="list-style-type: none"> <li>- Fitting a probabilistic graphical model</li> <li>- Testing hypotheses via data experiments</li> <li>- Applying a dimensionality reduction on a dataset to facilitate model training or gather insights</li> </ul>   | Data science<br>Algorithmic coding<br>Mathematics                     |
| Training deep learning models  | <ul style="list-style-type: none"> <li>- Using deep learning for a domain-specific application such as fraud detection, text summarization, machine translation, speech recognition, or object classification, detection, or segmentation</li> <li>- Tuning hyperparameters involved in neural network optimization</li> </ul>     | Deep learning<br>Algorithmic coding<br>Mathematics<br>Data science    |
| Accelerating training  | <ul style="list-style-type: none"> <li>- Setting up code to train a model on multiple machines in parallel</li> </ul>  | Domain-specific languages (for instance, CUDA)<br>Algorithmic coding  |
| Defining evaluation metrics (usually also involves a data product manager) | <ul style="list-style-type: none"> <li>- Choosing F1-score to evaluate a model's performance on a classification task</li> <li>- Implementing evaluation metrics such as accuracy, precision, recall, intersection over union, or mean average precision (mAP)</li> </ul>  | Machine learning<br>Algorithmic coding<br>Mathematics                 |



# Modeling work includes :

|  |  |   |
|--|--|---|
| Speeding up prediction time  | <ul style="list-style-type: none"><li>- Applying techniques such as pruning, quantization, or compression to reduce memory requirements</li><li>- Running inference speed vs. accuracy experiments on a model</li></ul>  | Machine learning<br>Algorithmic coding                      |
| Iterating over the virtuous cycle of machine learning projects: Idea, Code, Experiment | <ul style="list-style-type: none"><li>- Translating a business problem into a machine learning problem. For instance, depending on the quality and quantity of accessible data, an end-to-end network might lead to better results than a pipeline network</li><li>- Implementing the three-step cycle of ideating with your team, coding to set up experiments, analyzing results</li></ul> | Machine learning<br>Business acumen                         |
| Searching hyperparameters  | <ul style="list-style-type: none"><li>- Organizing experiments to get results in the shortest time period</li><li>- Setting up hyperparameter search experiments using tools such as AutoML</li></ul>  | Machine learning<br>Algorithmic coding                      |
| Keeping your knowledge up to date  | <ul style="list-style-type: none"><li>- Reading research papers</li><li>- Watching conference lectures or attending conferences</li></ul>  | Research<br>Mathematics<br>Data science<br>Machine learning |

# Deploy work includes :

| Subtask  | Examples   | Skills involved      |
|--|--|----------------------|
| Converting prototyped code into production code    | <ul style="list-style-type: none"><li>- Refactoring a repository's code</li><li>- Minimizing duplicate code</li><li>- Writing clean code to improve readability and consistency, for example, by following the PEP8 guidelines in Python</li></ul> | Software engineering |
| Setting up a cloud environment to deploy the model | <ul style="list-style-type: none"><li>- Mastering cloud tools and infrastructure provided by AWS, GCP, Azure, and the like</li><li>- Preparing files (usually model architecture and parameters) for deployment</li></ul>                          | Software engineering |
| Branching (version control)                        | <ul style="list-style-type: none"><li>- Designing a branching workflow, and using development, staging and production branches</li><li>- Participating in or leading code reviews</li></ul>  | Software engineering |
| Improving response times and saving bandwidth      | <ul style="list-style-type: none"><li>- Setting up load-balancing requirements with engineers in charge of AI Infrastructure</li></ul>   | Software engineering |

# Deploy work includes :

---

|  |   |  |
|--|---|--|
| Encrypting files that store model parameters, architecture, and data | <ul style="list-style-type: none"><li>- Understanding encryption at a high level and leveraging existing functions</li></ul>  | Software engineering                     |
| Building APIs for an application to use a model                      | <ul style="list-style-type: none"><li>- Setting up HTTP RESTful API services to facilitate communications between software components</li><li>- Setting up authorization and authentication to access the API</li></ul> | Software engineering                     |
| Retraining machine learning models (lifelong learning)               | <ul style="list-style-type: none"><li>- Monitoring changes in data distribution and staging model updates</li></ul>   | Software engineering<br>Machine learning |
| Fitting models on resource-constrained devices                       | <ul style="list-style-type: none"><li>- Pruning or quantizing a model so it fits memory requirements</li><li>- Deploying a model on a mobile device using TensorFlow</li></ul>  | Software engineering<br>Machine learning |

# Business Analysis work includes:

| Subtask  | Examples   | Skills involved   |
|--|--|---|
| Building data visualizations                             | <ul style="list-style-type: none"> <li>- Visualizing high-dimensional data in lower dimensions using methods such as PCA or t-SNE</li> <li>- Building and presenting graphs produced using Tableau, ggplot or matplotlib</li> <li>- Building visualizations in JavaScript, HTML, or CSS</li> </ul> | Domain-specific programming languages<br>Data science<br>Mathematics<br>Business acumen |
| Building dashboards for business intelligence            | <ul style="list-style-type: none"> <li>- Writing a script that periodically notifies business leaders of trends in the data</li> </ul>   | Domain-specific programming languages<br>Business acumen                                |
| Presenting technical work to clients or colleagues       | <ul style="list-style-type: none"> <li>- Preparing presentations (e.g., PowerPoint decks)</li> <li>- Communicating effectively with team members</li> <li>- Giving technical talks to present research outcomes</li> </ul>   | Communication<br>Business acumen  |
| Translating statistics into actionable business insights | <ul style="list-style-type: none"> <li>- Making marketing decisions based on analysis of various sources</li> </ul>  | Data Science<br>Business Acumen   |

# Business Analysis work includes:

---

|   |   |   |
|---|---|---|
| Analyzing datasets                              | <ul style="list-style-type: none"><li>- Plotting a correlation matrix to analyze covariates</li><li>- Computing statistical variables such as mean, variance, and mode</li><li>- Segmenting customers into groups</li></ul>   | Data science<br>Algorithmic coding<br>Mathematics     |
| Running experiments to evaluate deployed models | <ul style="list-style-type: none"><li>- Working with the deployment team to evaluate business performance of a deployed model</li><li>- Helping the deployment team make decisions</li><li>- Translating model performance into business outcomes such as revenue</li></ul> | Data science<br>Algorithmic coding                    |
| Running A/B tests                               | <ul style="list-style-type: none"><li>- Optimizing web pages</li><li>- Evaluating systems in production</li></ul>   | Data science<br>Algorithmic coding<br>Business acumen |

# AI infrastructure work includes:

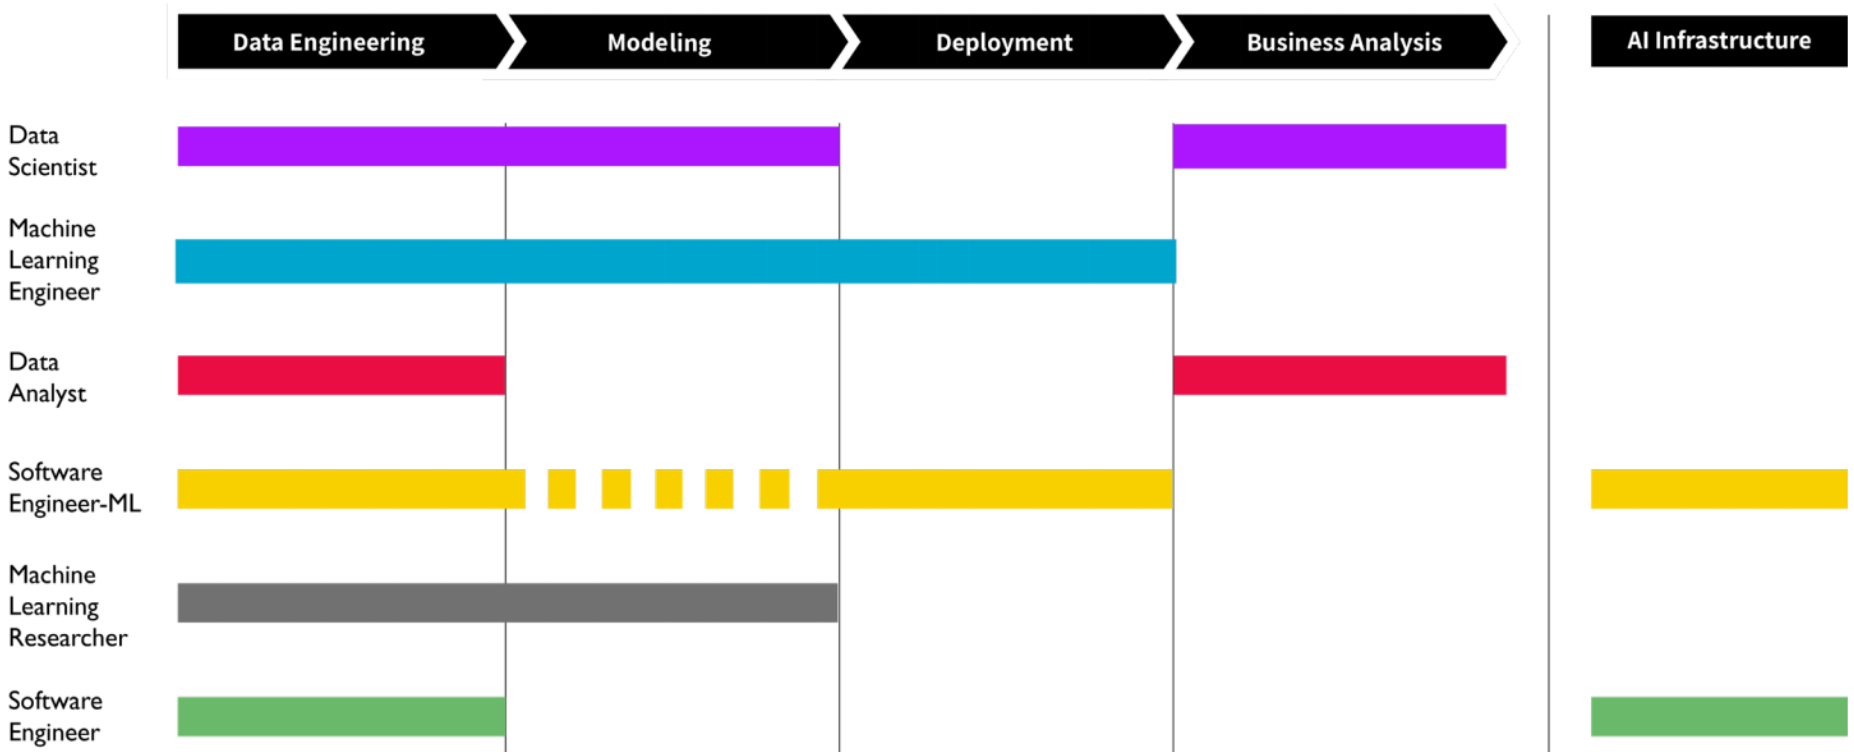
| Subtask  | Examples  | Skills involved                                   |
|--|---|---|
| Making software design decisions                   | - Reducing latency by locating a model close to data  | Software engineering                              |
| Building distributed storage and data-base systems | - Building databases (SQL, NoSQL, MySQL, Cassandra, etc.) to store data and facilitating access by other team members | Software engineering<br>Domain-specific languages |
| Designing for scale                                | - Adding GPU compute or storage as needed   | Software engineering                              |
| Maintaining software infrastructure                | - Managing software upgrades and driving stability through automated monitoring and alerting                          | Software engineering                              |
| Networking   | - Controlling access to all infrastructure elements   | Software engineering                              |



# AI infrastructure work includes:

---

|                                     |  |                      |
|-------------------------------------|--|----------------------|
| Networking                          | - Controlling access to all infrastructure elements  | Software engineering |
| Securing data and models            | - Building security features that allow for production deployments into regulated organizations, satisfying the needs for privacy and security | Software engineering |
| Writing tests                       | - Writing unit and functional tests for multiple components across tasks of the AI project life cycle  | Software engineering |
| Carrying out various software tasks | - Building a labeling program, A/B testing framework, or analysis environment  | Software engineering |



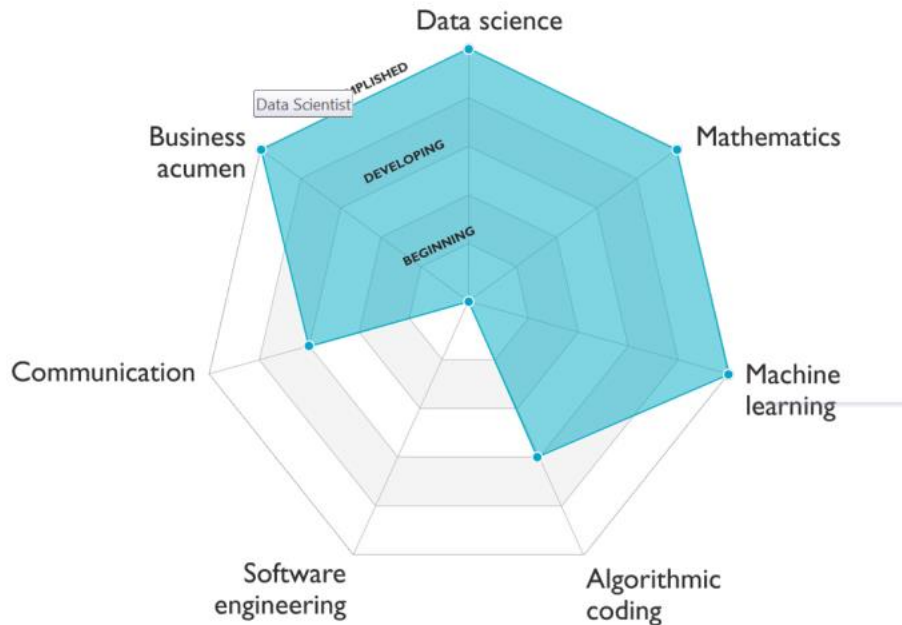
# Data Scientist

---

- **Data science organizations** help a firm's leaders make scientific or data-driven decisions to run their business more effectively. Team members collect data, analyze datasets, and suggest hypotheses and actions.
- **Machine learning organizations** automate tasks to reduce costs or scale products. The output is the automation itself achieved by collecting data, training models, and deploying them.

# Data Scientist

## SKILLS



## TOOLS

- **Modeling** in Python using packages such as numpy, scikit-learn, TensorFlow, and PyTorch
- **Data engineering** in Python and/or SQL or other domain-specific query languages
- **Business analysis** in Python, R, other domain-specific tools such as Tableau or Excel, or presentation software applications such as PowerPoint or Keynote
- **Collaboration and workflow** using a version control system such as Git, Subversion, or Mercurial along with a command line interface (CLI) such as Unix and an integrated development environment (IDE) such as Jupyter Notebook or Sublime


# Machine Learning Engineer vs Data Scientist

## Machine Learning engineer:

- Knowledge of Python programming, Keras, PyTorch etc.
- Data wrangling and APIs
- Supervised and unsupervised algorithms
- Statistics for AI
- Natural Language Processing
- Deep learning and computer vision

## Data scientist:

- Statistical skills
- Programming languages
- Data visualisation
- Big data platforms
- Machine learning



## Data Engineer

## Data Scientist

*Design, build, and arrange  
data to be cleaned.*

*Create hypothesis, test,  
analyze, and translate  
using clean data*

Advanced Programming

Advanced Math/Statistics

Distributed Systems

Analysis

Machine Learning

Programming

Data Pipelines

Big Data






Advanced Analytics



# Data Analyst



## Data Scientist vs Data Analyst

| Features   | Data Scientist   | Data Analyst   |
|--|--|--|
|  Background   | A Data Scientist deals with various data operations.                         | A Data Analyst's role is related to data cleaning, transforming and generating inferences from data. |
|  Scope        | Involved with several underlying data procedures                             | Involvement is limited to small data and static inferences.  |
|  Type of Data | Handles structured & unstructured data                                       | Deals with structured data only  |
|  Skills      | Possesses knowledge of mathematics, statistics & machine learning algorithms | Has problem solving skills, knowledge of basic statistics  |
|  Tools      | Proficient in SAS, Python, R, TensorFlow, Hadoop, Spark                      | Knows Excel, SQL, R (in some cases), Tableau   |

# Data Analyst

---



Data Engineer

I am responsible for preparing the foundation that data scientist use for their work



Data Scientist

I analyze data to identify patterns and trends to deliver recommendations for future




Data Analyst

I analyze data to summarize the past in visual form

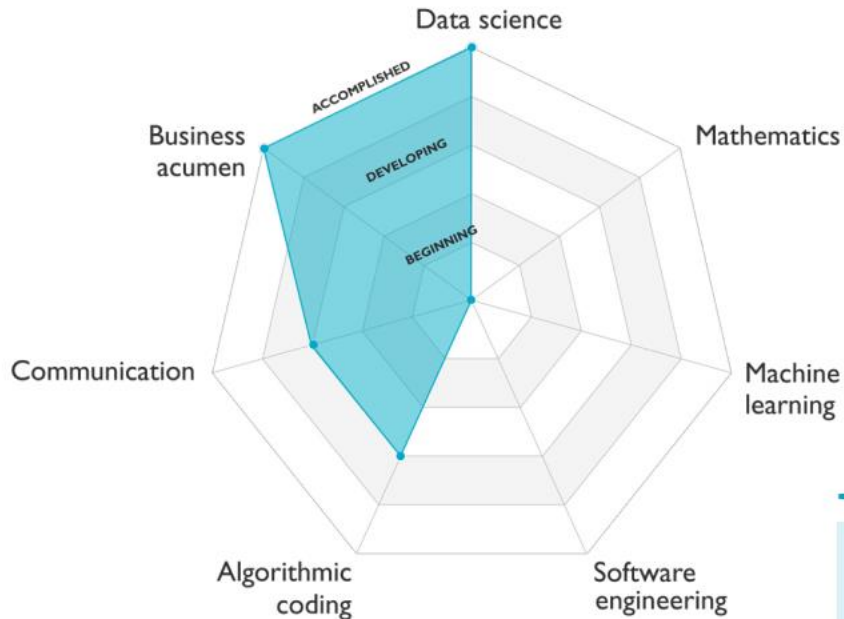
# Data Analyst

---



# Data Analyst

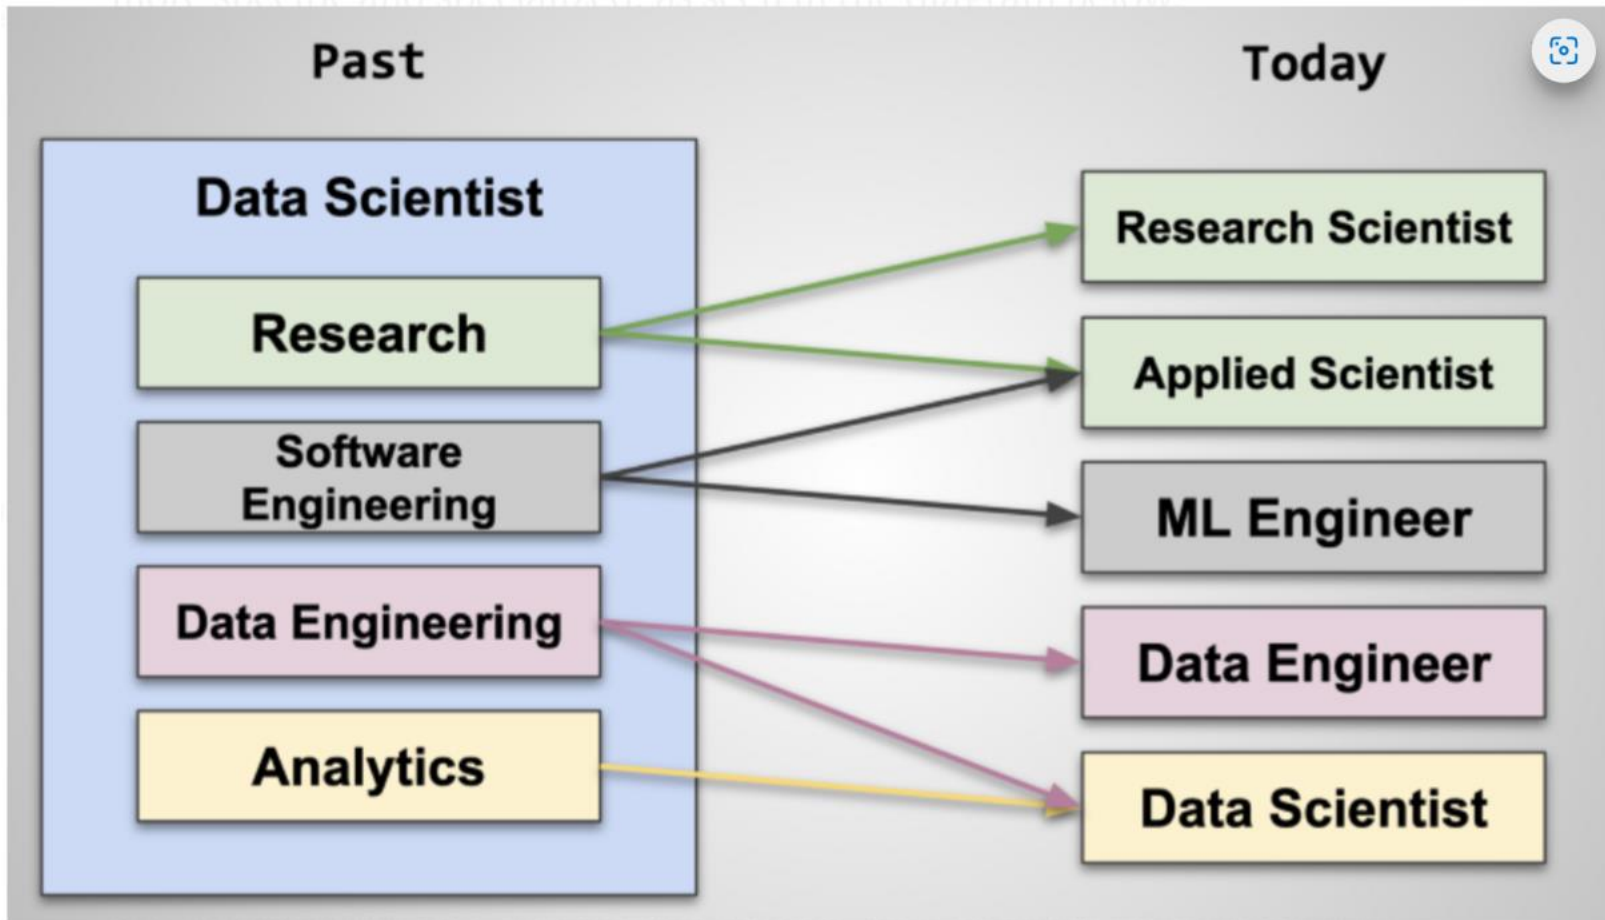
## SKILLS



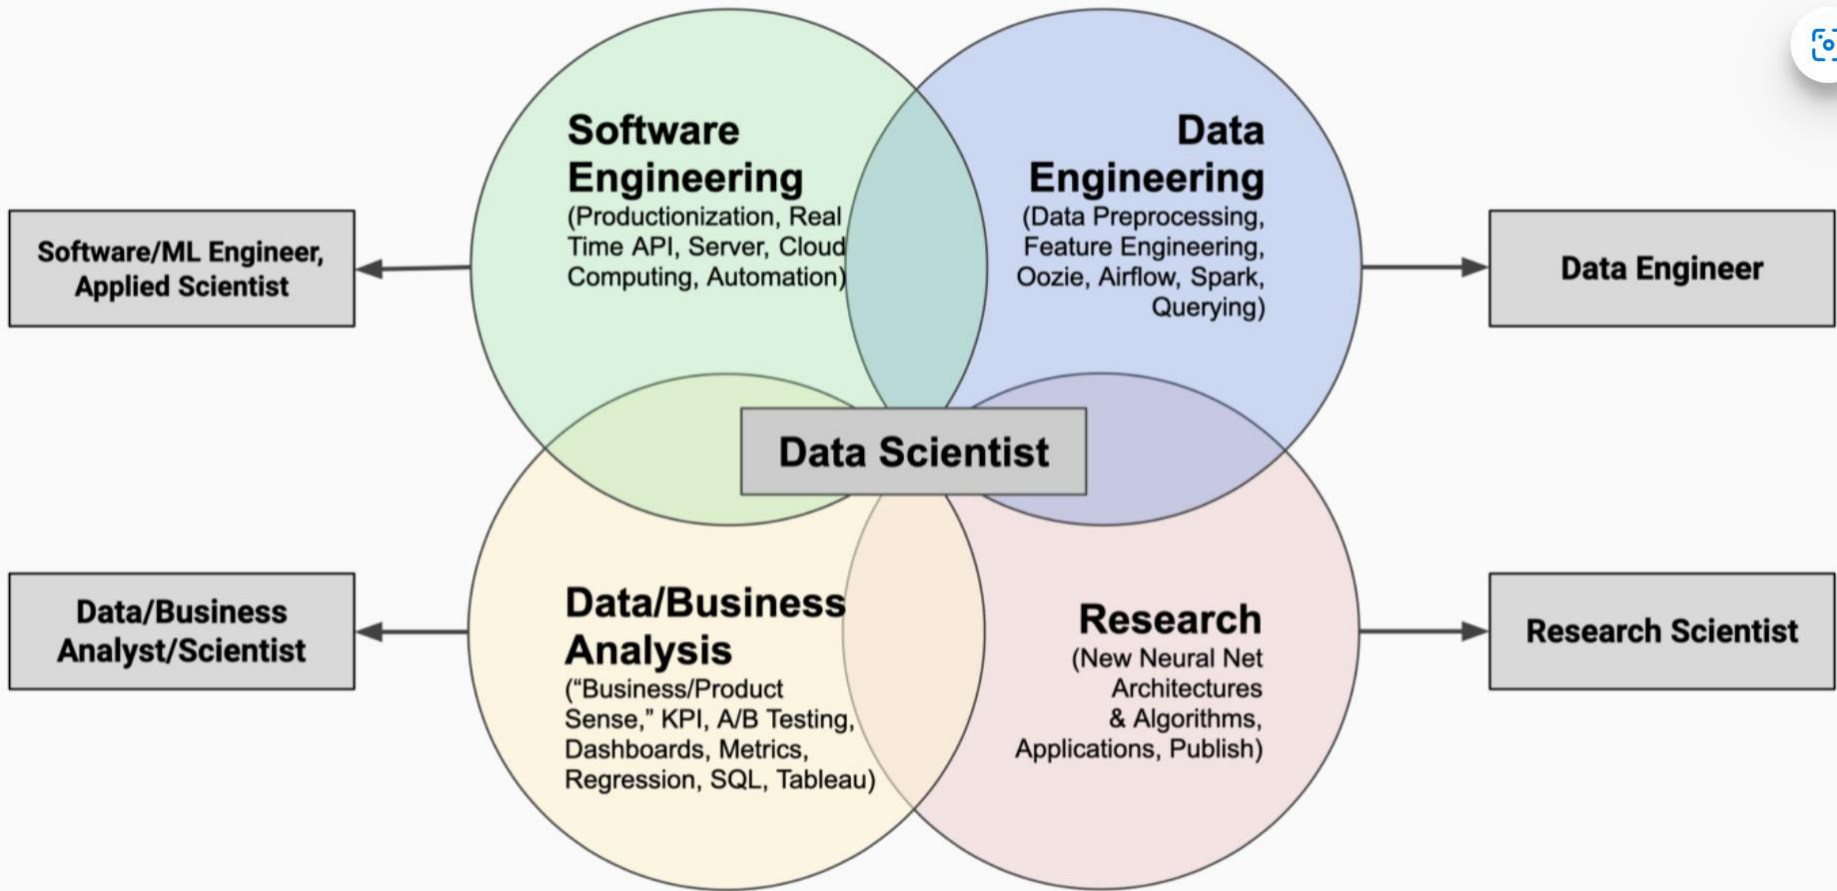
## TOOLS

- **Data engineering** in Python and/or SQL or other domain-specific query languages
- **Business analysis** in Python, R, other domain-specific tools such as Tableau and Excel, presentation software applications such as PowerPoint and Keynote, and external software services for A/B testing

# Machine learning Engineer



# Machine learning Engineer






# Machine Learning Engineer vs Data Scientist

## Machine Learning engineer:

- Knowledge of Python programming, Keras, PyTorch etc.
- Data wrangling and APIs
- Supervised and unsupervised algorithms
- Statistics for AI
- Natural Language Processing
- Deep learning and computer vision

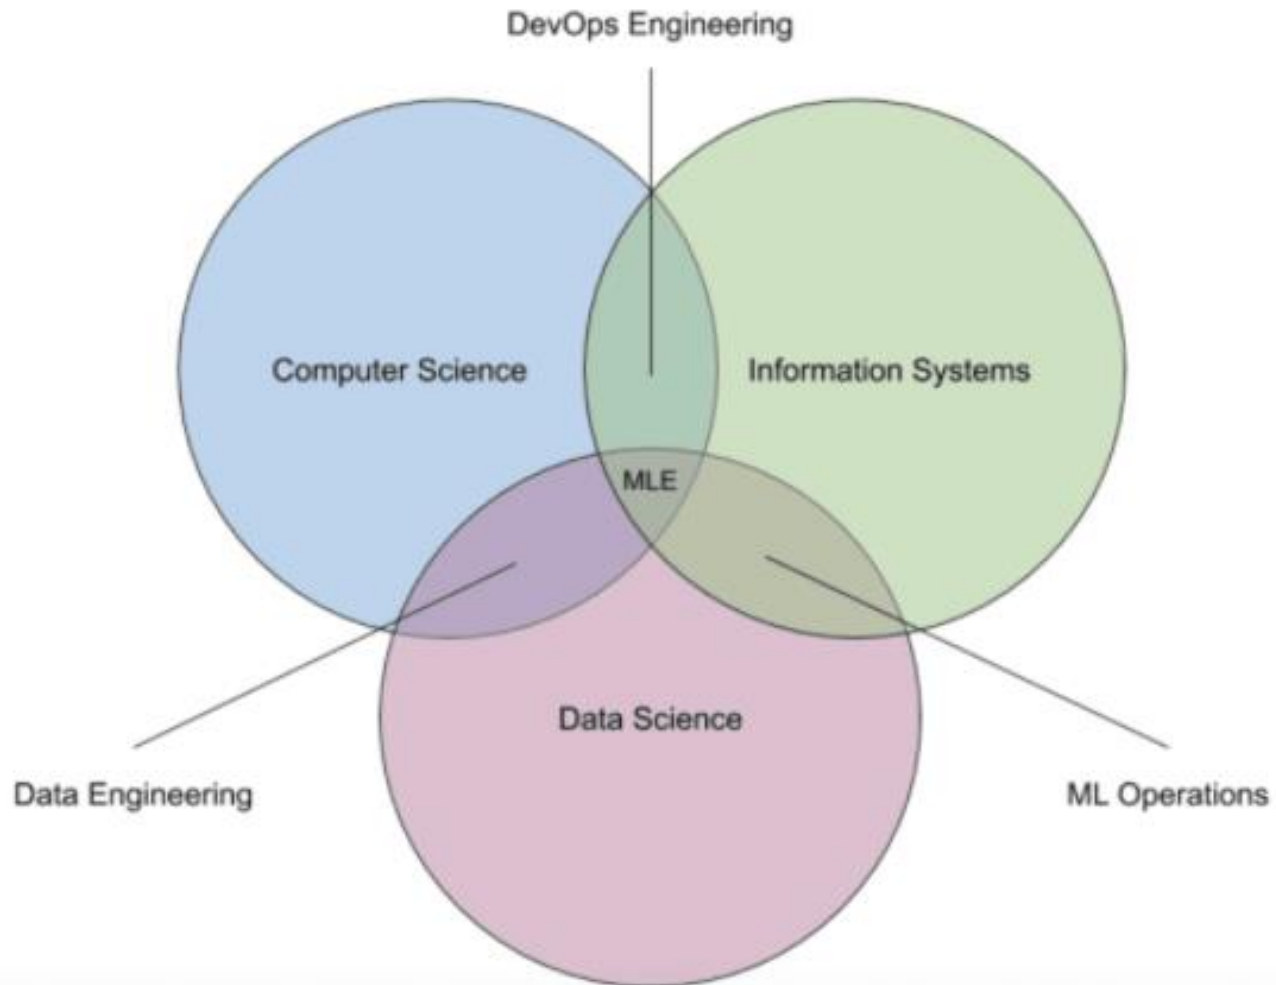
## Data scientist:

- Statistical skills
- Programming languages
- Data visualisation
- Big data platforms
- Machine learning




# Machine learning Engineer

---




# Machine learning Engineer

## SKILLS



## TOOLS

- **Data Engineering** in Python and/or SQL or other domain-specific query languages
- **Modeling** in Python using packages such as numpy, scikit-learn, TensorFlow, and PyTorch
- **Deployment** using an object-oriented programming language (such as Python, and Java, C++,) and cloud technologies such as AWS, GCP, and Azure
- **Collaboration and workflow** using a version control system (for instance, Git, Subversion, and Mercurial), a command line interface (CLI) like Unix, an integrated development environment (IDE) such as Jupyter Notebook, and Sublime, and an issue tracking product like JIRA



## **Software Development Team**

Focused on developing a feature, product component, or an entire project.

Codes and then delivers the web or app developed software to the client.

**VS**



## **Support & Operations Team**

The support team ensures the software development team product's functionality over time.


Is in charge of incident management, product testing, troubleshooting, bug fixing, monitoring and alerting, etc.

**Cafeto**

# What is DevOps?



Developers & Testers

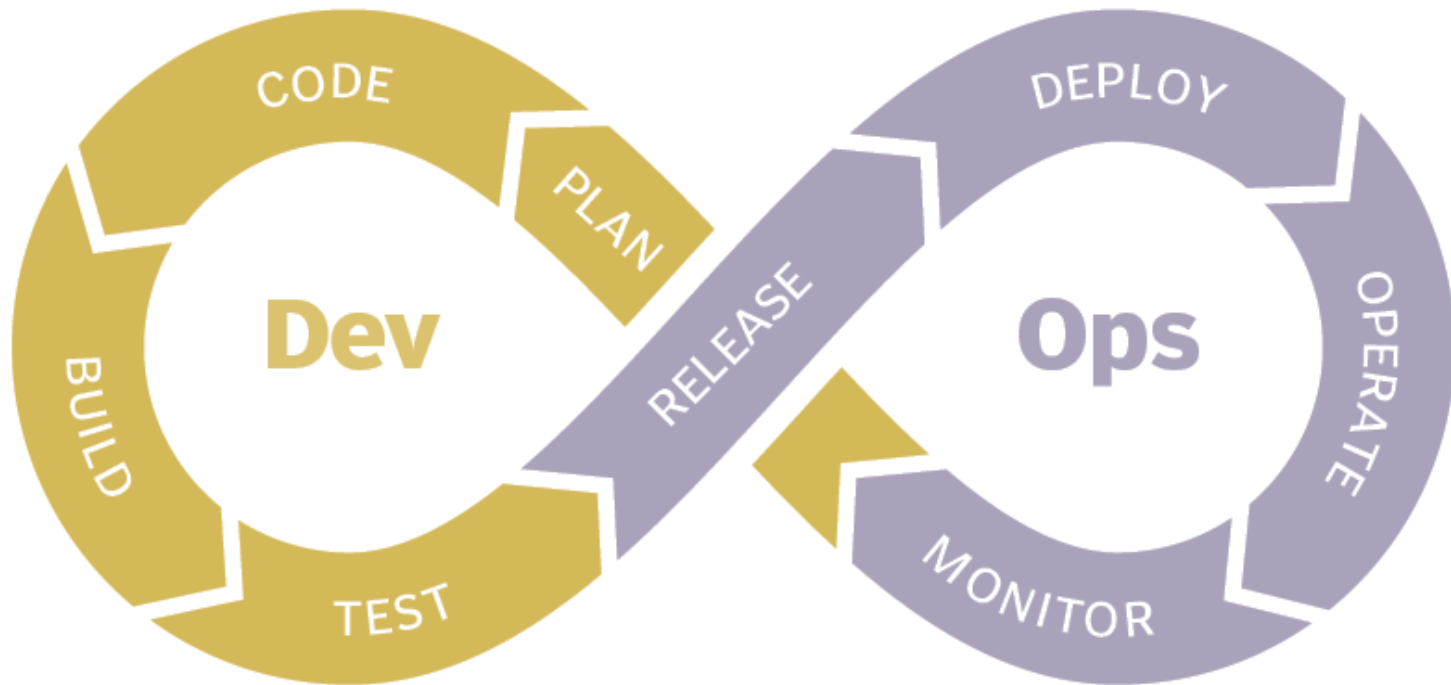


IT Operations





# DevOps infinity loop



# From **Devops** to **MLOPS**

**MLOps = ML + DEV + OPS**




## **Develop**

- Algorithm Training + Testing
- ETL (Data Pipelines)
- Continuous Integtation / Continuous Deployment

## **Operate**

- Continuous Delivery
- Model Inference
- Monitoring and Management



# Machine learning ReSearcher

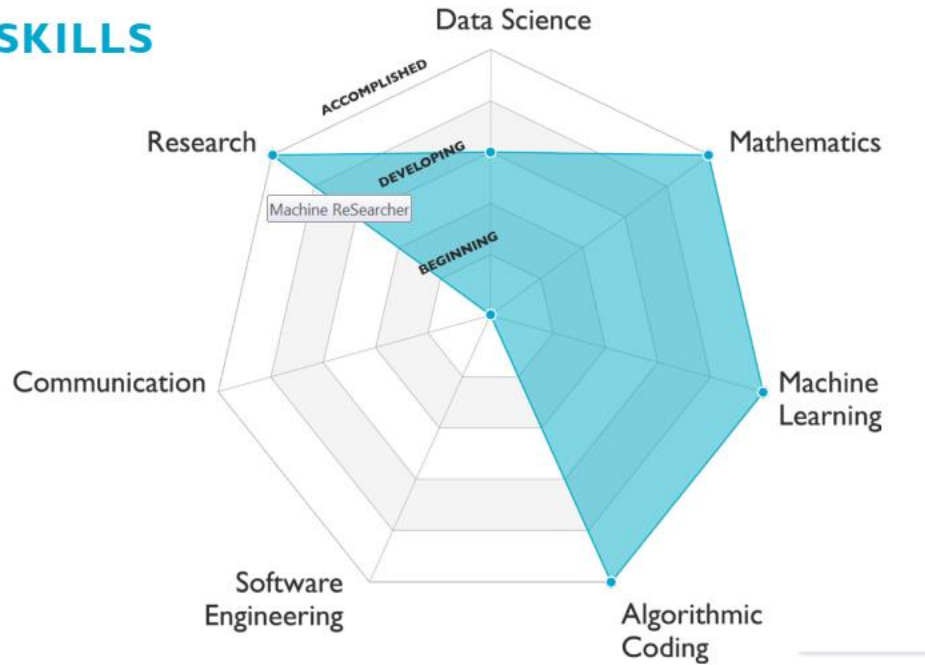
---

Machine learning research is really all about the *science*. A machine learning researcher is trying to push the boundaries of science, specifically in the field of Artificial Intelligence. These people typically have a Masters or PhD in CS and have many publications in top machine learning conferences. They're super popular in the research space!



# Machine learning ReSearcher

## SKILLS

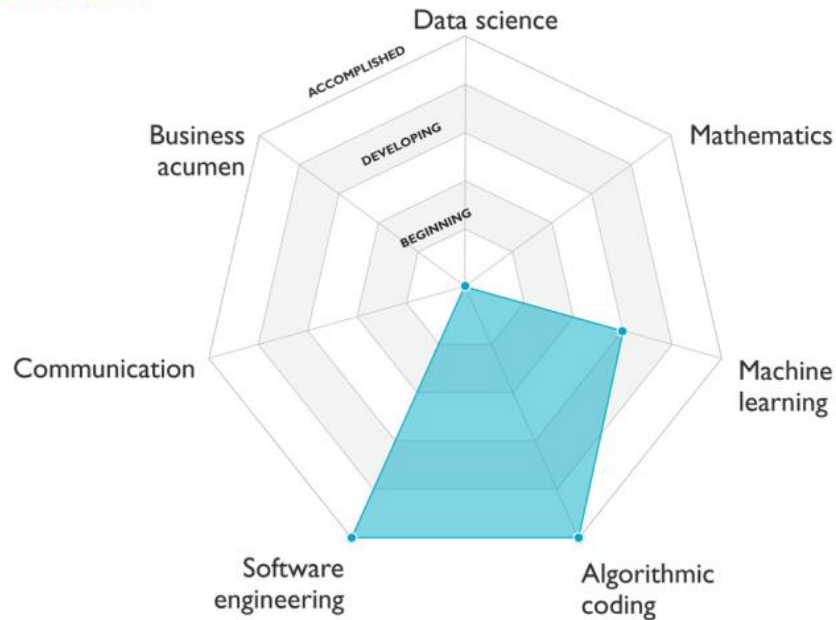


## TOOLS

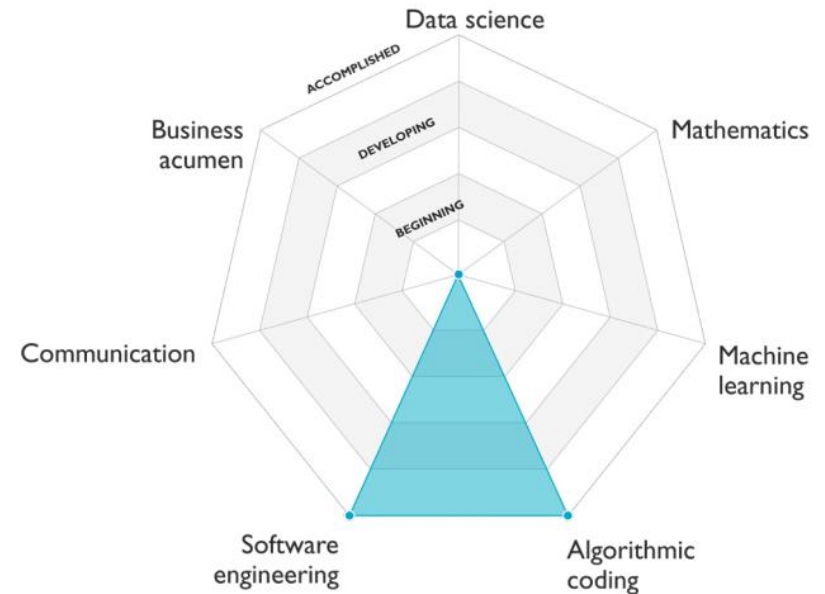
- **Data Engineering** in Python and/or SQL (or other domain-specific query languages)
- **Modeling** in Python using packages such as numpy, scikit-learn, TensorFlow, PyTorch, and the like
- **Collaboration and workflow** using a version control system like Git, Subversion, and Mercurial, a command line interface (CLI) like Unix, an integrated development environment (IDE) such as Jupyter Notebook or Sublime, and an issue tracking product like JIRA
- **Research** by following updates via channels such as Twitter, Reddit, Arxiv, and conferences such as NeurIPS, ICLR, ICML, CVPR, and ACM

# Software Engineer -Machine Learning VS Software Engineer

## SKILLS



## SKILLS



# Software Engineer -Machine Learning VS Software Engineer

---

## TOOLS

- **Modeling** in Python using packages such as numpy, scikit-learn, TensorFlow, and PyTorch
- **Data engineering** in Python and/or SQL (or other domain-specific query languages)
- **Deployment and AI infrastructure** using an object-oriented programming language such as Python, Java, or C++ and cloud technologies such as AWS, GCP, or Azure
- **Collaboration and workflow** using a version control system like Git, Subversion, or Mercurial, a command line interface (CLI) like Unix, an integrated development environment (IDE) like Jupyter Notebook or Sublime, and an issue tracking product such as JIRA

## TOOLS

- **Data engineering** in Python and/or SQL (or other domain-specific query languages)
- **AI infrastructure** using an object-oriented programming language such as Python, Java, or C++ and cloud technologies such as AWS, GCP, and Azure
- **Collaboration and workflow** using a version control system like Git, Subversion, or Mercurial, a command line interface (CLI) like Unix, an integrated development environment (IDE) like Jupyter Notebook and Sublime, and an issue tracking product like JIRA



